



# Node embeddings

# We can define each node by its connections

Each node is represented by a vector



You can use it to predict something about the node:

- but that would mean thousands or millions of parameters!
- and only provides information at the local level

# Node embeddings

Idea: Create a low(er)-dimension representation of the node  
→ **Similar nodes** should have a similar representation

What do we mean with similar nodes?

- Nodes with the same neighborhood (useful to predict e.g. clustering or assortative attributes)
- Nodes with the same role (useful to predict e.g. structural equivalence or disassortative attributes)
- Nodes with the same metadata

# Unsupervised vs supervised learning

**Supervised learning:** Output is available. Performance = discrepancy between predicted output and real output


- Regression
- Classification (e.g. link prediction!)

**Unsupervised learning:** No labels/output. Performance = reduction of some error

- Clustering (e.g. cluster points so they are as close as possible within clusters, as far as possible between clusters)
- Dimensionality reduction (e.g. combine variables to maximize the amount of variability explained)

# Option 1: Spectral methods

Related to characteristic eigenvectors of matrices associated with the network



## Principal Component Analysis (PCA)


- Correlated variables  $\rightarrow$  Linear combination of orthogonal variables
- Eigenvectors corresponding to the largest  $k$  eigenvectors of the covariance matrix



## Singular Value Decomposition (SVD)

- Eigenvectors corresponding to the largest  $k$  eigenvectors of  $A^T A$

## Laplacian Eigenmaps (LE)


- Assumes that the nodes lie on a low-dimensional (with some constraints)
- Tries to find a mapping that minimizes the distances between all nodes
- That mapping is created by the eigenvectors corresponding to the smallest  $k$  eigenvectors of the normalized Laplacian matrix  $D^{-1/2}(D-A)D^{-1/2}$








# Option 2: Matrix factorization

## Singular Value Decomposition (SVD)



## Non-negative Matrix Factorization (NMF)







# Option 3: Shallow Neural Networks

In text analysis and image recognition we have a lot of neural network tools because the data is regular

- Text analysis: Chain (nodes = words)
- Images: Lattices (nodes = pixels)

## **Word2vec** (SkipGram/CBOW)

Distributional hypothesis: similar words will be surrounded by similar words (you will know a word by the company it keeps)

What words appear around “Network” in text?

- Network Science
- Network Analysis

→ Words *science* and *analysis* are similar

# Word2Vec

## Step 1: Create co-occurrence matrix

- I like deep learning
- I like NLP
- I enjoy flying

counts	I	like	enjoy	deep	learning	NLP	flying	.
I	0	2	1	0	0	0	0	0
like	2	0	0	1	0	1	0	0
enjoy	1	0	0	0	0	0	1	0
deep	0	1	0	0	1	0	0	0
learning	0	0	0	1	0	0	0	1
NLP	0	1	0	0	0	0	0	1
flying	0	0	1	0	0	0	0	1
.	0	0	0	0	1	1	1	0


# Word2Vec

## Step 2: Train the model

- *Positive examples*: word  $i$  used to predict word  $j$  (context word in the co-occurrence matrix)
- *Negative examples*: word  $i$  used as a negative example to predict word  $k$  (not a context word)

Two vectors are similar if they have a high dot product ( $\sim$ cosine similarity)

- Vector associated to "deep":  $w1["\text{deep}",:]$
- Vector associated to "learning":  $w2[:,\text{"learning"}]$




Intuition:

- Modify  $W1$  and  $W2$  so target embeddings are close (have a high dot product) to context embeddings for nearby words and further from context embeddings for noise words that don't occur nearby

Speech and language processing: an introduction to natural language processing, computational linguistics, and speech recognition (Jurafsky and Dan, 2009)

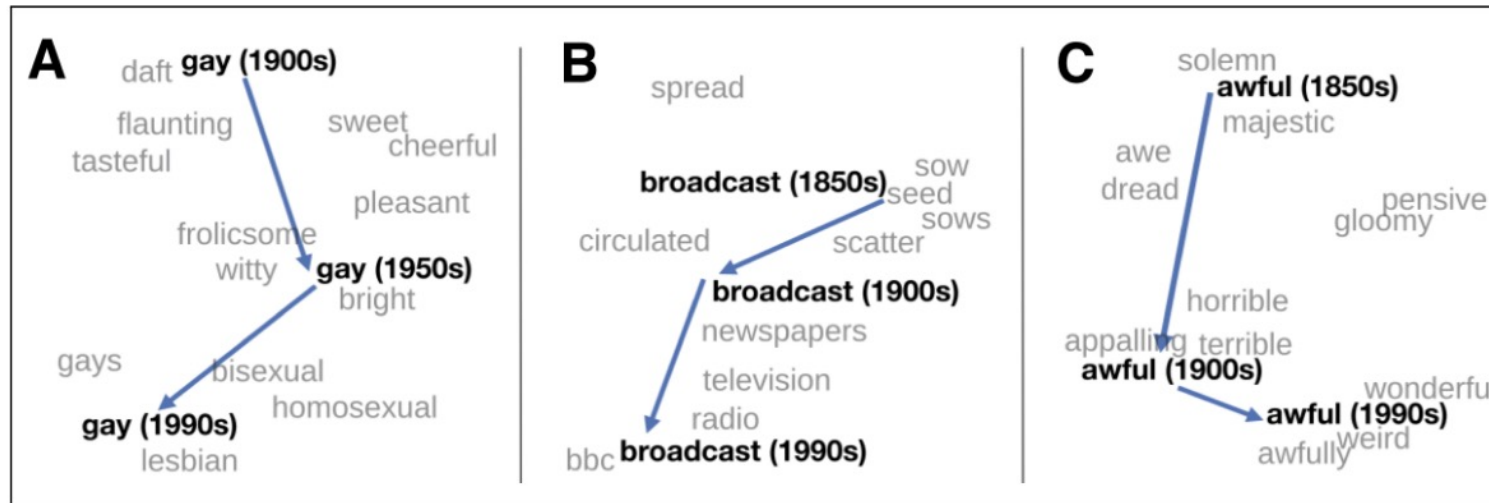
# Word2Vec



**Figure 6.15** The parallelogram model for analogy problems ([Rumelhart and Abrahamson, 1973](#)): the location of  $\overrightarrow{\text{vine}}$  can be found by subtracting  $\overrightarrow{\text{apple}}$  from  $\overrightarrow{\text{tree}}$  and adding  $\overrightarrow{\text{grape}}$ .

# Option 3: Shallow Neural Networks

## Word2vec

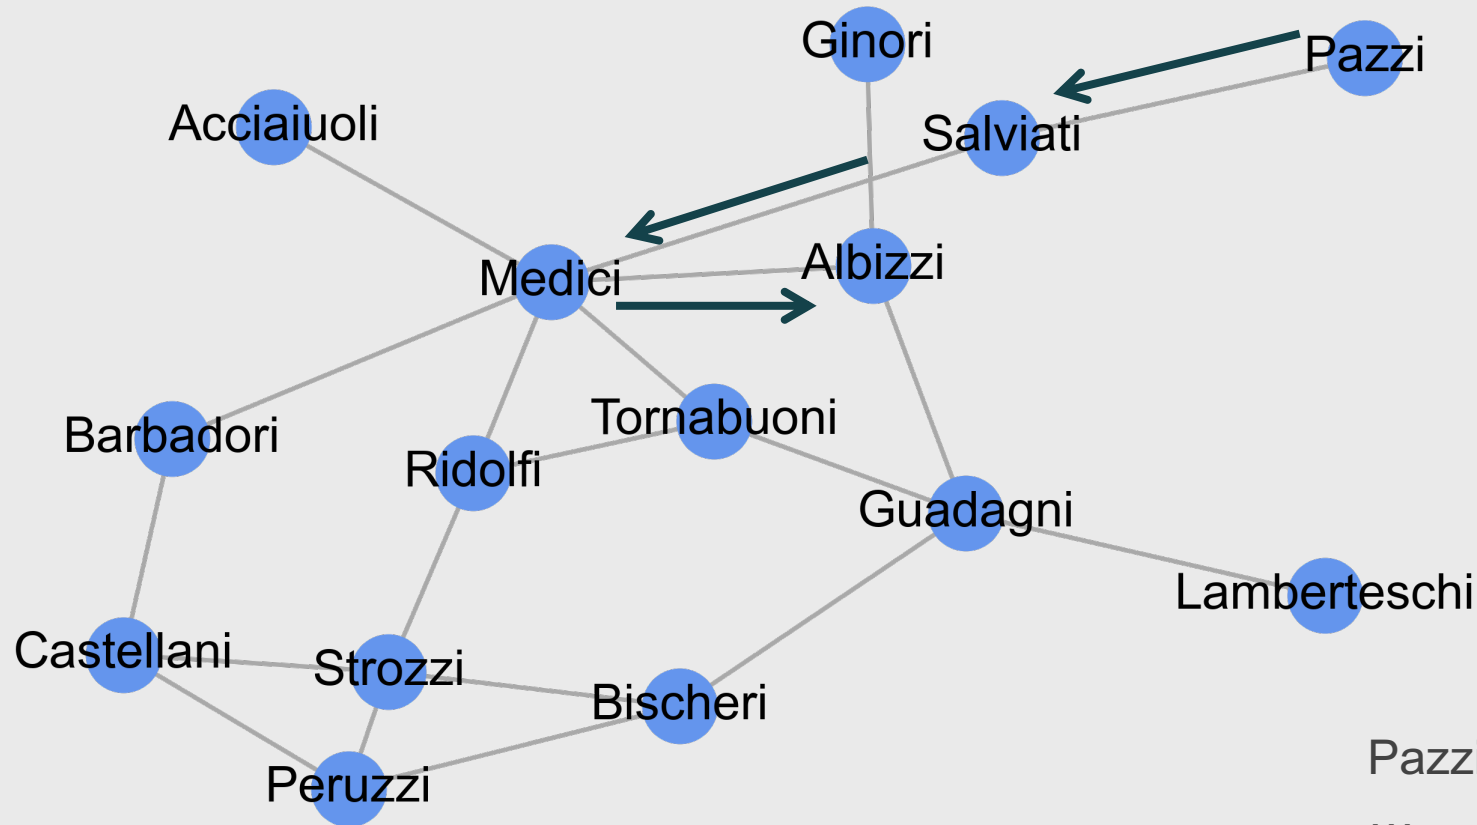


**Figure 6.17** A t-SNE visualization of the semantic change of 3 words in English using word2vec vectors. The modern sense of each word, and the grey context words, are computed from the most recent (modern) time-point embedding space. Earlier points are computed from earlier historical embedding spaces. The visualizations show the changes in the word *gay* from meanings related to “cheerful” or “frolicsome” to referring to homosexuality, the development of the modern “transmission” sense of *broadcast* from its original sense of sowing seeds, and the pejoration of the word *awful* as it shifted from meaning “full of awe” to meaning “terrible or appalling” (Hamilton et al., 2016b).

# In networks: node2vec (deepwalk)

Generate "sentences" using random walks.

The more times two nodes appear in the same random walk, the more similar their embeddings will be



Pazzi -> Salviati -> Medici -> Albizzi

...

# Node2Vec (deepwalk)

## Step 1: Create co-occurrence matrix

Pazzi -> Salviati -> Medici -> Albizzi

Medici -> Albizzi -> Guadagni -> Medici

...	Pazzi	Salviati	Guadagni	Medici	Albizzi
Pazzi		1		1	1
Salviati	1				
Guadagni					
Medici	1		1	1	1
Albizzi	1			1	

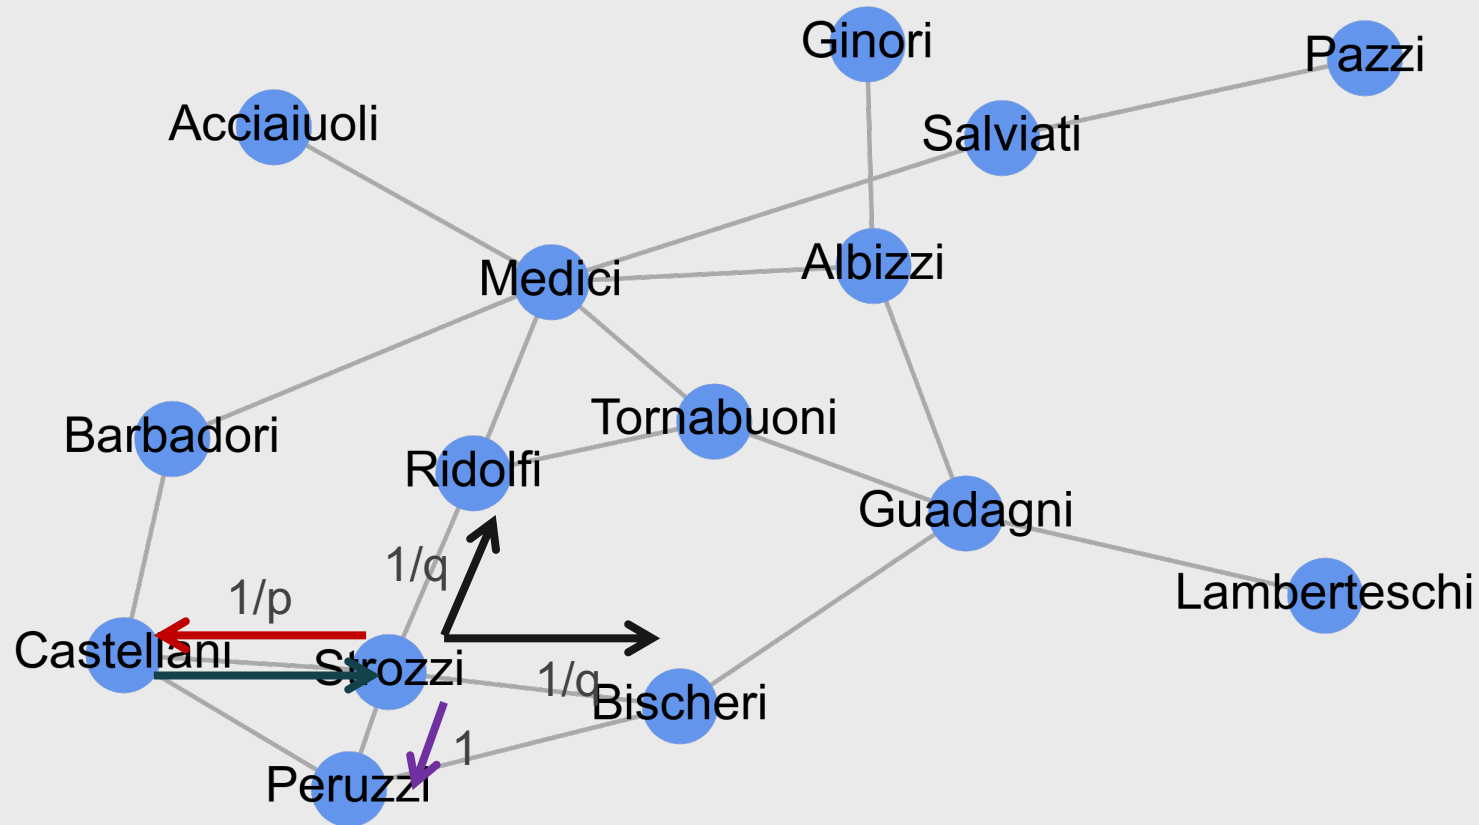
**Step 2: Create embeddings** representing how similar the neighbors of each node are

# node2vec

Difference with deepwalk: generate “sentences” using biased random walks.

Q = controls probability of going to new nodes

P = controls probability of going back to previous node





# Using node2vec (deepwalk)



$q = 0.5$

~ similarity reflecting clusters



$q = 2$

~ similarity reflecting structural roles

Figure 3: Complementary visualizations of Les Misérables coappearance network generated by *node2vec* with label colors reflecting homophily (top) and structural equivalence (bottom).

# What to do with the embeddings

## 1) Node classification

$X$  = Embedding

$Y$  = whatever we want to predict

## 2) Link prediction

$X$  = combination of the two embeddings (e.g. dot product, or  $\text{np.abs}(v1-v2)$ )


$Y$  = link/no link

# Option 4: Deep neural networks

We created the embeddings in an unsupervised (or self-supervised) way. But we can do it in a supervised way, so **node embeddings are similar if they are related to the same outcome**.

Many methods (GCN, GAT, graphSAGE...)


- **Intuition:** Nodes aggregate information from their neighbors using neural networks




# Option 4: Deep neural networks

- **Intuition:** Network neighborhood defines a computation graph

Every node defines a computation graph based on its neighborhood!



INPUT GRAPH



# Primer on machine learning

# Machine learning

"A computer program is said **to learn from experience  $E$**  with respect to some class of **tasks  $T$**  and **performance measure  $P$**  if its performance at task  $T$ , as measured by  $P$ , improves with experience  $E$ ." (Samuel/Mitchell, 1959)

- Experience: Data
- Task: Goal
- Performance measure: Accuracy,  $R^2$ , etc

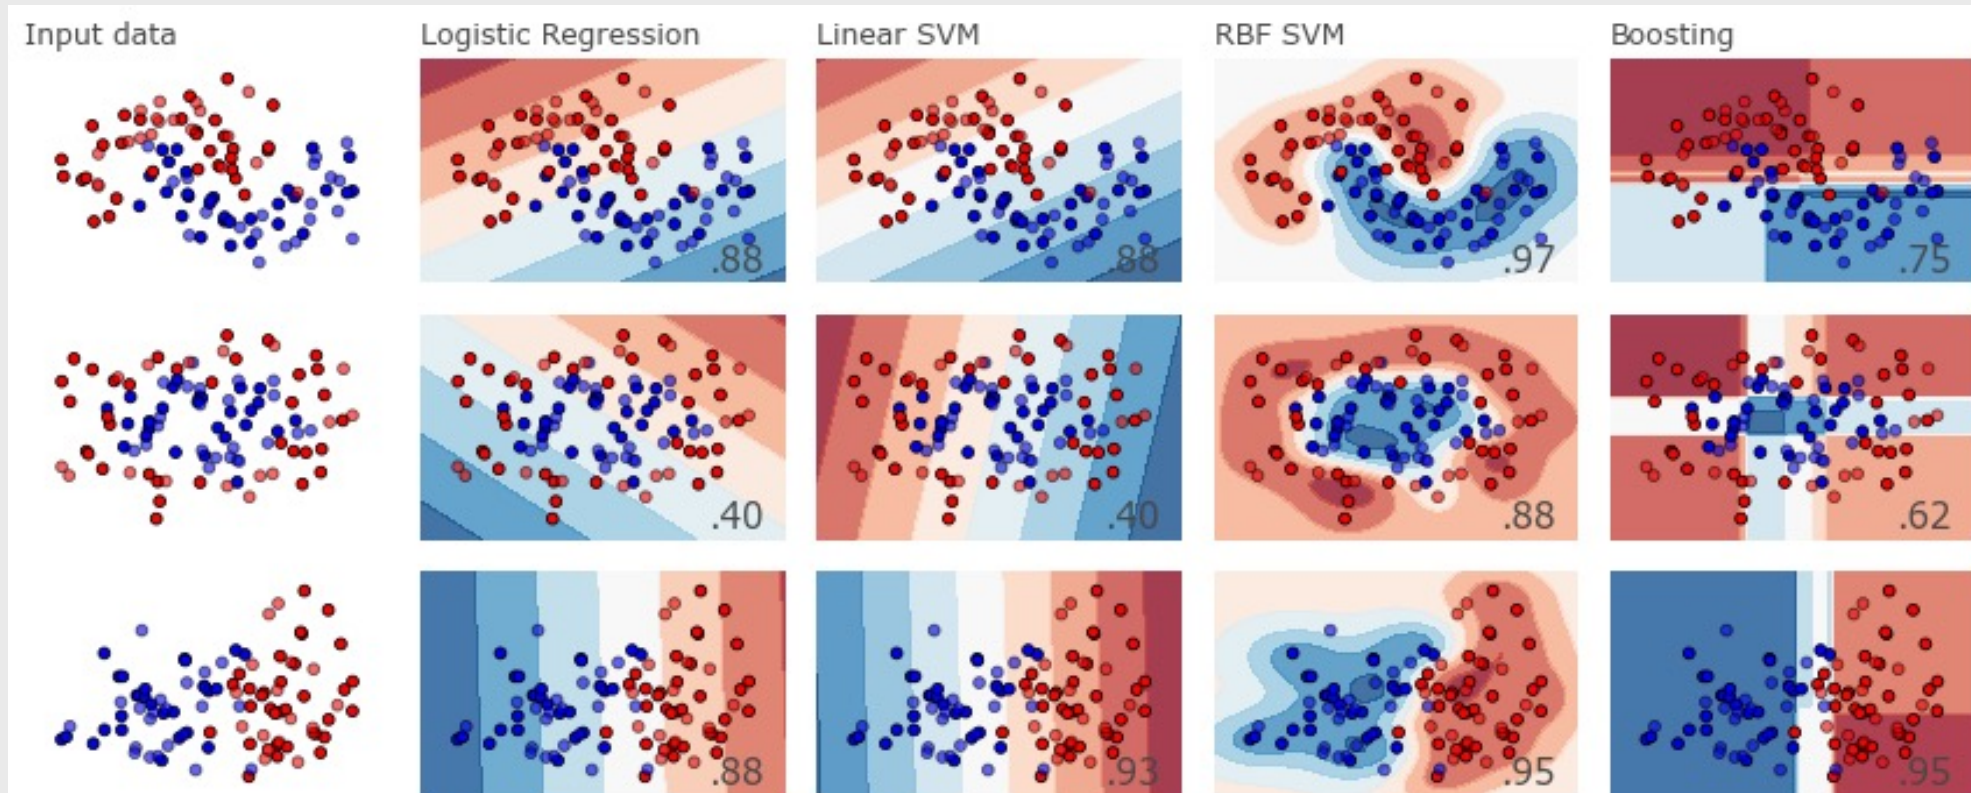
# Machine learning

- Typically focuses on **large, high-dimensional** datasets with **complex interactions** between features
- **Output-driven:**
  - Typically aims to solve a problem (rather than to test a hypothesis): e.g. link prediction
  - Emphasizes predictive accuracy: Uses theory (to build new features) if it improves accuracy
  - The model works as intended, on new data

*Cannot and should not replace thinking about causation*

# One use in link prediction: stacking classifiers

- No best algorithm: *Stacking models for nearly optimal link prediction in complex networks; Ghasemian, Galstyan, Airoldi, Clauset (2020)*
- Each algorithm gives you a score: we can combine them in a model
- Methods for this: (Penalized) logistic regression; Support Vector Machines; Boosting





# Main issues in Machine Learning


## Issue 1: Overfitting

- Lots of data and features → can be a recipe for disaster
- Evaluation of overfitting: cross-validation
- Prevention of overfitting: Regularization, weak learners, dropout, etc

## Issue 2: Interpretability of the model


- Complex models are more difficult to interpret
- New measures of interpretability

# Issue 1: Overfitting



Adding more variables always decreases  $R^2$  (within-sample prediction error). How to estimate real prediction error?

- One option: adjust for the degrees of freedom
- The “ML” option: evaluate prediction accuracy in an external dataset



# 1: Evaluate overfitting using a validation dataset




**Training set**

**Validation**



**Training dataset** → Use to train different models

**Validation dataset** → Evaluate out-of-sample prediction error

(Test dataset) → Evaluate out-of-sample prediction error of final model



# 1: Evaluate overfitting using a validation dataset




But with this:

- We reduce the training dataset (number of observations)
- We validate on a small dataset (maybe not representative)


Solution → Cross-validation

# 1: Evaluate overfitting using cross-validation



Do you want to understand the error due to the splitting? → Run this procedure several times with random splittings


# 1: Hyperparameter tuning



**Hyperparameter tuning using cross-validation** →  
Balance between flexibility and overfitting:

- **Regularization** (e.g. sum of  $|\text{coefs}| < I$ )
- Ensembles:
  - Train trees with different data
    - Bootstrap
    - Subset of predictors
  - Use shallow trees
- Neural networks:
  - Train disabling neurons (dropout)
- Early stopping

## Issue 2: Interpretability



If you can't  
explain it simply,  
you don't  
understand it well  
enough.

ALBERT EINSTEIN

## Typical workflow

*# Split data into training and testing to evaluate generalization error (more on this later)*

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
```

*# Hyperparameter tuning using your model: e.g., linear\_model.LogisticRegression(), svm.SVC(), ensemble.HistGradientBoostingClassifier()*

```
parameters = {'kernel':('linear', 'rbf'), 'C':[1, 10]} # (e.g. for a SVC)
```

```
clf = GridSearchCV(model() , parameters)
```

```
clf.fit(X_train, y_train)
```

```
print(clf.best_params_)
```


*# Use the best model to predict the labels (link/no link) in the testing data*

```
y_pred = clf.best_estimator_.transform(X_test)
```

```
print(F1_score(y_test, y_pred))
```



# Structure of the challenge



**Protein-protein interaction network in *S. cerevisiae***

Clustering  $\sim 0$ ; Assortativity  $\sim -0.2$

We have removed some edges, your objective is to predict those accurately.

### **We give you:**

- Graph: Used for training
- Test dataset (a series of node pairs, some with a link associated)

### **How:**

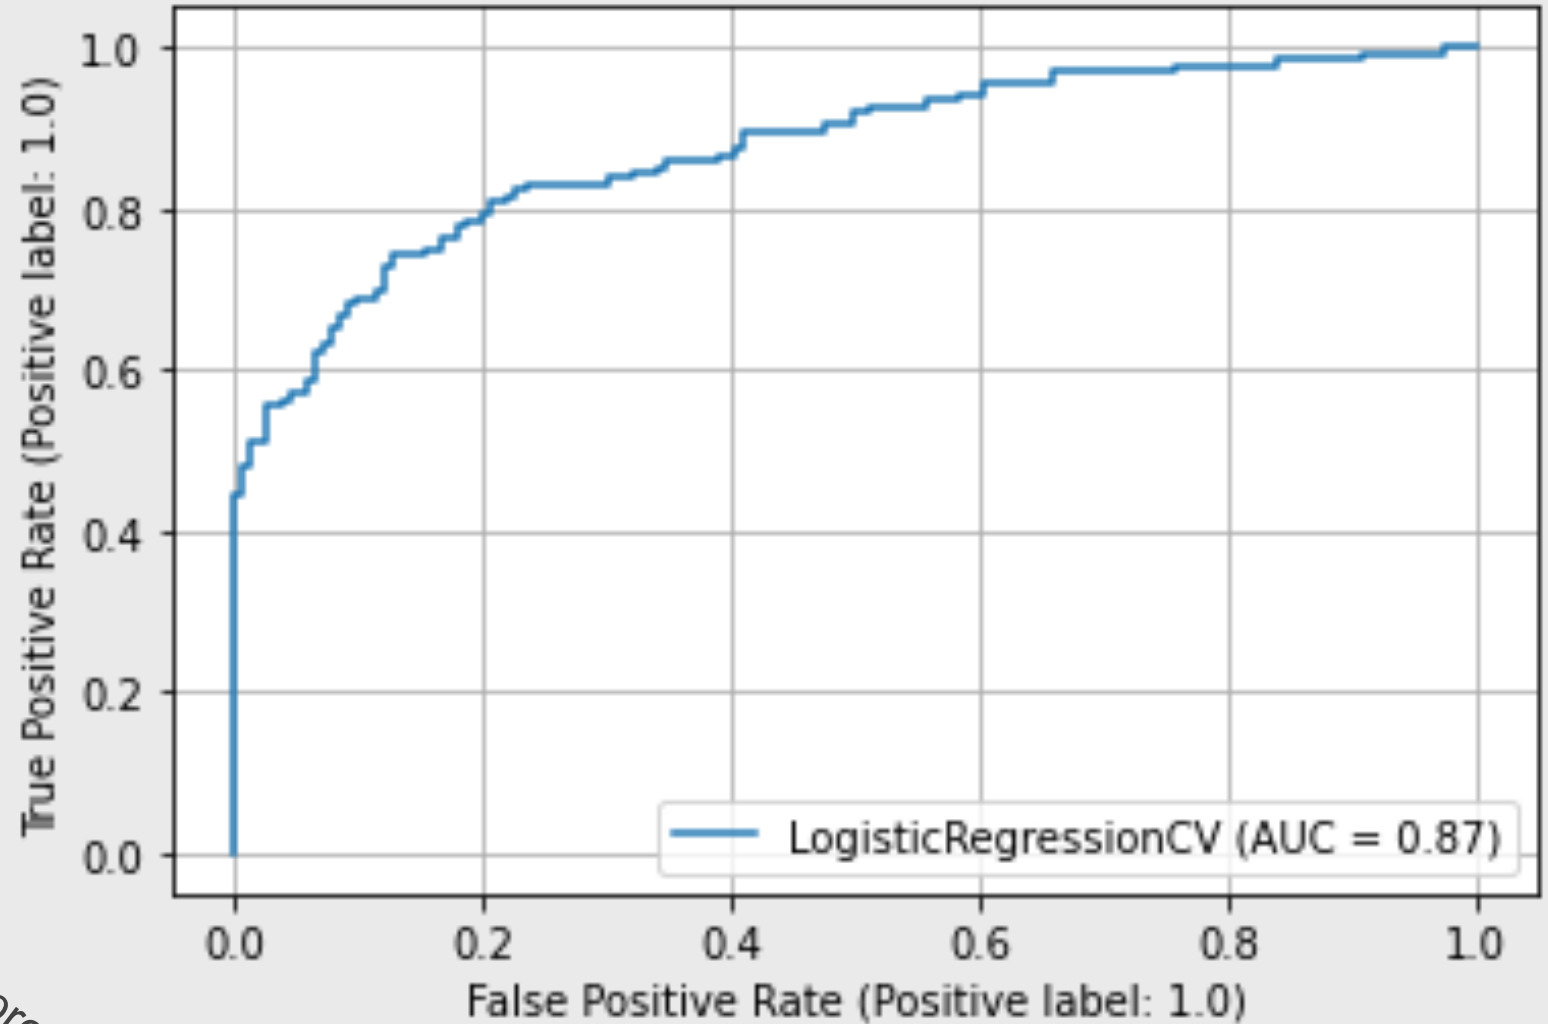
- Methods based on similarity
- Methods based on paths
- SBM
- Methods based on embeddings
  - Spectral methods
  - Matrix factorization
  - Node2vec

# Evaluation

(all links  
predicted  
correctly)

$$\begin{aligned} \text{TPR (sensitivity)} &= \\ &= \frac{\text{Links predicted correctly}}{\text{Node pairs with a link}} \end{aligned}$$

(no links predicted)



(all no links predicted as links)

$$\text{FPR} = \frac{\text{No links predicted as links}}{\text{Node pairs with no link}}$$